

# Flow Matching Corrected Posterior Estimation (FMCPE)

Pierre-Louis Ruhlmann

Pedro L. Rodrigues, Michael Arbel, Florence Forbes

EDF workshop on Metamodels

# About me

- PhD student at Inria Grenoble (3rd year)
- Statify/Thoth team
- Research focus: Simulation-based inference
- Supervisors: Pedro L. Rodrigues, Michael Arbel, Florence Forbes

*Inria*



Statify

# About me



Inria



Statify

# Introduction

# Statistical Inference

Given observations  $x$  and unknown parameters  $\theta$ , we use **Bayes' rule**:

$$p(\theta|x) = \frac{p(x|\theta)p(\theta)}{p(x)}$$

- $p(\theta|x)$  - **Posterior**: what we want to compute
- $p(x|\theta)$  - **Likelihood**: probability of data given parameters
- $p(\theta)$  - **Prior**: initial beliefs about parameters
- $p(x)$  - **Evidence**: normalizing constant

# Statistical Inference

$$p(\theta|x) = \frac{p(x|\theta)p(\theta)}{p(x)}$$

## Bayesian Inference

**Input:** Model likelihood  $p(x|\theta)$

**Methods:** MCMC, Variational Inference

**Output:** Posterior  $p(\theta|x)$

## Simulation-Based Inference

**Input:** Simulator  $S : \theta \mapsto x$

**Methods:** *(See next section)*

**Output:** Posterior  $p(\theta|x)$

**Key difference:** SBI works when likelihood is intractable

# Approximate Bayesian Computation (ABC)

## Algorithm

**Inputs:** Simulator  $S$ ,  $\mathbf{y}_{\text{obs}}$ , prior  $p(\theta)$ , distance  $\rho$ , threshold  $\epsilon$

**For**  $i = 1, \dots, n$ :

1. Sample  $\theta_i \sim p(\theta)$
2. Simulate  $\mathbf{x}_i = S(\theta_i)$
3. Keep  $\theta_i$  if  $\rho(\mathbf{x}_i, \mathbf{y}_{\text{obs}}) < \epsilon$

## Python Code

```
# Inputs
S = simulator
y_obs = observed_data
prior = p_theta
rho = distance_function
epsilon = threshold

samples = []
for i in range(n):
    # 1. Sample from prior
    theta_i = prior.sample()
    # 2. Simulate
    x_i = S(theta_i)
    # 3. Keep if distance < epsilon
    if rho(x_i, y_obs) < epsilon:
        samples.append(theta_i)
```

## Limitations:

- Poor scaling with dimension
- Requires summary statistics
- Not amortized

**References:** Beaumont et al. (2002), Sisson et al. (2018)

# Neural Posterior Estimation (NPE)

## Algorithm

**Inputs:** Simulator  $S$ , budget  $N_{\text{train}}$ ,  $N_{\text{epochs}}$ , neural approximator  $q_\phi$ , summary net  $h_\omega$

**Generate training data:**

Sample  $\{(\theta_i, \mathbf{x}_i)\}_{i=1}^{N_{\text{train}}}$

**For**  $e = 1, \dots, N_{\text{epochs}}$ :

Minimize KL divergence:

$$\mathcal{L}(\phi, \omega) = -\mathbb{E}[\log q_\phi(\theta | h_\omega(\mathbf{x}))]$$

## Python Code

```
# Inputs
S = simulator
N_train = budget
N_epochs = num_epochs
q_phi = neural_approximator
h_psi = summary_net

# Generate training data
thetas = prior.sample(N_train)
xs = S(thetas)

# Training loop
for epoch in range(N_epochs):
    # Minimize KL divergence
    loss = -q_phi.log_prob(thetas, h_psi(xs)).mean()
    loss.backward()
    optimizer.step()
```

**Benefits:** Amortized inference, no hand-crafted statistics needed

**Limitations:** Prone to overfitting, sensitive to distribution shift

**References:** Papamakarios & Murray (2016), Greenberg et al. (2019)

# Challenges in SBI

- **Expensive simulations** - High computational cost
- **Exploit the simulator** - In certain cases where the simulator is not a black box, we can leverage information (e.g. its gradient) to improve inference.
- **Model misspecification** - Main focus of this work.

# Model Misspecification

# Types of Misspecification and Objective

## Types of misspecification:

- **Simulator misspecification:** Model assumptions are wrong or incomplete (e.g., missing physics)
- **Prior misspecification:** Prior distribution does not reflect true parameter ranges
- **Observation noise:** Noise model in likelihood is misspecified

## Approaches to address misspecification:

- **Robustness:** Ensemble methods or conservative posteriors
- **Detection:** Identify presence of misspecification
- **Correction:** Build new posterior accounting for discrepancy

→ In this work, we focus on correcting simulator misspecification.

# Misspecification of the simulator

The data generating process  $\mathbf{y} \sim p^*$

$$p^* \notin \{p(\cdot|\theta) \mid \theta \in \Theta\}$$

## Example: Damped Pendulum

**High-fidelity model (ground truth):**

$$\mathbf{y}(t) = e^{-\alpha t} A \cos(\omega_0 t) + \epsilon$$

Parameters:  $\theta = (\omega_0, A)$ , damping  $\alpha \sim \mathcal{U}(0, 1)$

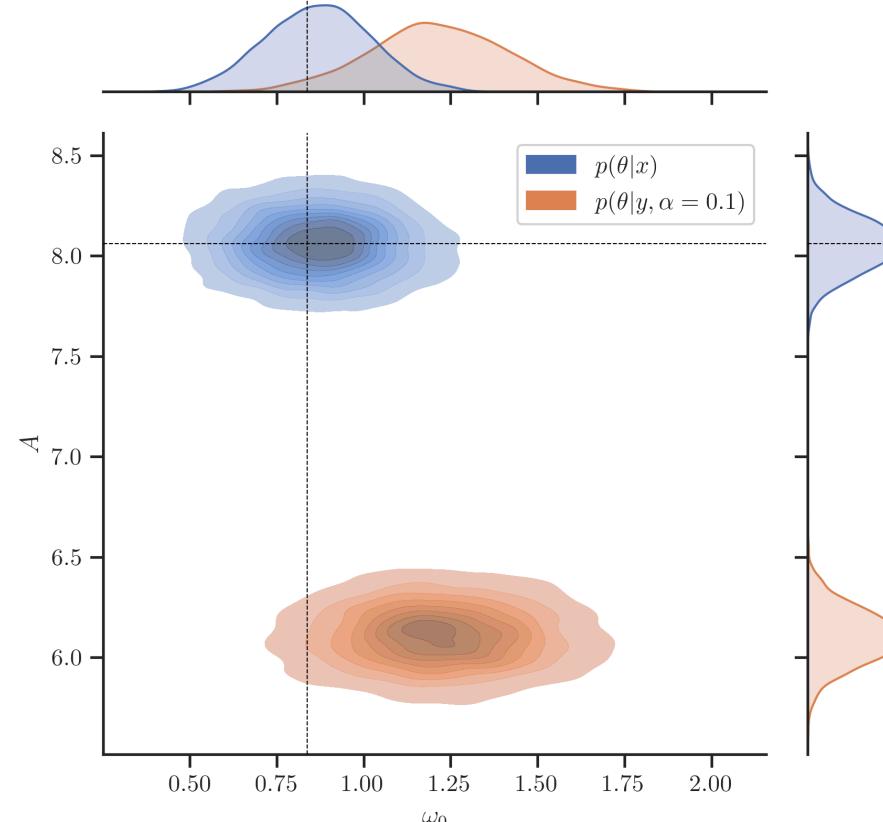
**Low-fidelity simulator (misspecified):**

$$\mathbf{x}(t) = A \cos(\omega_0 t) + \epsilon$$

**Missing damping term!  $\alpha = 0$**

→ Even a small amount of misspecification can lead to highly biased posteriors !

NPE:  $p(\theta|x)$  vs  $p(\theta|y, \alpha = 0.1)$



# Data setting

There are two main ways to approach model misspecification in SBI:

## Case 1

### 1. Simulation data:

Access to simulated pairs  $\{\theta_i, \mathbf{x}_i\}_{1:m}$

### 2. Observations:

Independent observations  $\{\mathbf{y}_j\}_{1:n}$   
( $\mathbf{y}$  does not depend on  $\theta$ )

### 3. Examples:

(Add examples here)

## Case 2

### 1. Simulation data:

Access to simulated pairs  $\{\theta_i, \mathbf{x}_i\}_{1:m}$

### 2. Calibration set:

Few paired observations  $\{\theta_j, \mathbf{y}_j\}_{1:n}$   
(*Multi-fidelity simulation, expensive ground truth*)

### 3. Examples:

(Add examples here)

→ In this work, we will focus on Case 2.

# Existing Methods

## MF-NPE

A first approach to correct the misspecification is to fine-tune the posterior estimator on the calibration set.

Assuming we have data from a low-fidelity simulator  $(\theta, \mathbf{x}) \sim p(\theta)p(\mathbf{x} \mid \theta)$ , we train a neural density estimator using:

$$\mathcal{L}_{NPE}(\phi) = \mathbb{E}_{\theta, \mathbf{x}} [-\log q_\phi(\theta | \mathbf{x})]$$

Then, using our small but precise calibration set  $(\theta_i, \mathbf{y}_i) \sim p(\theta, \mathbf{y})$ , we fine-tune the posterior estimator using the same loss:

$$\mathcal{L}_{NPE}(\phi) = \mathbb{E}_{\theta, \mathbf{y}} [-\log q_\phi(\theta | \mathbf{y})]$$

**Reference:** Krouglova et al. (2025). Multifidelity Simulation-based Inference for Computationally Expensive Simulators. *arXiv:2502.08416*.

# ROPE: Transporting $\mathbf{y}$ to $\mathbf{x}$

The ROPE method proposes to use optimal transport to match  $p^*(\mathbf{y})$  and  $p(\mathbf{x})$ .

**Input:** Simulator  $S(\theta, \epsilon)$ , calibration data  $\{(\theta^i, \mathbf{y}^i)\}_{i=1}^{N_c}$  and test set  $\{\mathbf{y}^i\}_{i=1}^{N_o}$

**Step 1:** Train n.n.  $h_\omega$  and density estimator  $q_\phi$

$$\mathcal{L}_{NPE}(\omega, \phi) = \mathbb{E}_{\theta, \mathbf{x}} [-\log q_\phi(\theta | h_\omega(\mathbf{x}))]$$

**Step 2:** Fine-tune  $h_\omega$  on calibration data  $\{(\theta^i, \mathbf{y}^i)\}$  (Initialize  $g = h$ )

$$\mathcal{L}(\varphi) = \sum_{i=1}^{N_c} \|g_\varphi(\mathbf{y}^i) - \mathbb{E}_\epsilon[h_\omega(S(\theta^i, \epsilon))]\|^2$$

**Step 3:** Solve OT problem between  $\mathbf{x}^j$  and  $\mathbf{y}^i$  with cost:

$$C(\mathbf{x}^j, \mathbf{y}^i) = \|h_\omega(\mathbf{x}^j) - g_\psi(\mathbf{y}^i)\|^2$$

**Output:** Corrected posterior  $\tilde{p}(\theta | \mathbf{y}) = \sum_j^{N_s} N_o P_{ij}^* q_\phi(\theta | h_\omega(\mathbf{x}^j))$

## Key takeaways of ROPE

- Two main components: The alignment of summary statistics using fine tuning and the transport between  $\mathbf{x}$  and  $\mathbf{y}$
- The correspondence between  $\theta_i$  and  $\mathbf{y}^i$  is only used to define the OT cost
- The "test" set is used for the computation of the optimal transport. The method can't be applied to new data
- The hypothesis  $\mathbf{y} \perp \theta \mid \mathbf{x}$  is required to compute the posterior

# Our Approach

Flow Matching Corrected Posterior Estimation

# Main Idea: Learn to Correct the Posterior

## The Challenge:

- We have a posterior  $q_\phi(\theta|\mathbf{x})$  trained on cheap simulations  $\mathbf{x}$
- We have a small **calibration set**  $\mathcal{D}_{\text{cal}} = (\theta_i, \mathbf{y}_i)$  from the true process
- Goal: Estimate  $p(\theta|\mathbf{y})$  by leveraging the calibration set and  $q_\phi(\theta|\mathbf{x})$

## Two step approach:

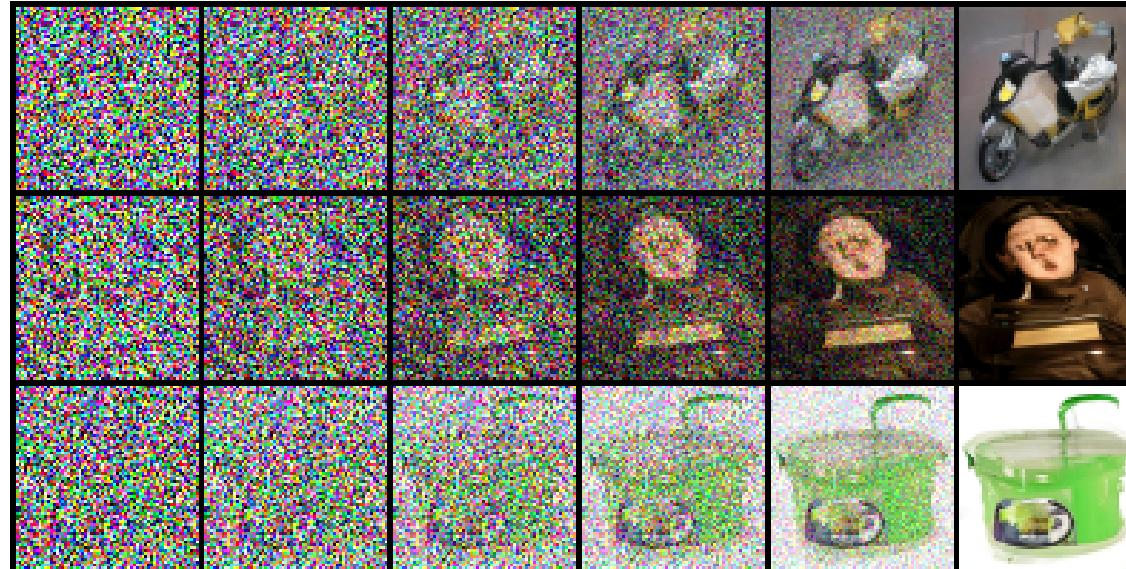
1. Design a **source distribution**  $\pi(\theta|\mathbf{y})$  which acts as a first guess for the target  $p(\theta|\mathbf{y})$
2. Use **Flow Matching** to move samples from the source  $\pi(\theta|\mathbf{y})$  to  $p(\theta|\mathbf{y})$

## Key Advantages:

- Leverages abundant simulations AND expensive calibration data efficiently
- No independence assumption  $\mathbf{y} \perp \theta|\mathbf{x}$  required (as opposed to RoPE).
- Can be plugged-in on top of any SBI methods that provides  $\pi(\theta|\mathbf{x})$

# Flow Matching: Key Idea

Flow Matching learns to transport samples between distributions



## Three key elements:

1. **Base distribution** we can sample from (e.g., Gaussian)  $\rightarrow q_0(x)$
2. **Target distribution** (arbitrary, potentially complex)  $\rightarrow q_1(x)$
3. **Transport mechanism** to move samples  $q_0(x)$  from  $q_1(x) \rightarrow p_t(x)$

## References:

Lipman et al. (2023). Flow Matching for Generative Modeling. *ICLR'23*.

Liu et al. (2023). Flow Straight and Fast: Learning to Generate and Transfer Data with Rectified Flow. *ICLR'23*.

# Flow Matching

**Objective:** Match 2 distributions  $q_0, q_1$  via an ODE, i.e. learn a probability path  $p_t(\mathbf{x})$  with  $p_0 = q_0$  and  $p_1 = q_1$

$$\frac{d\psi_t(\mathbf{x})}{dt} = u_t(\psi_t(\mathbf{x})) \quad \psi_0(\mathbf{x}) = \mathbf{x}$$

**Relationship between flow and probability path:**

$$\mathbf{x}_t = \psi_t(\mathbf{x}_0) \text{ with } \mathbf{x}_0 \sim q_0 \quad \Rightarrow \quad \mathbf{x}_t \sim p_t$$

**Boundary conditions:**

$$\psi_0 = \text{Id} \quad \text{and} \quad \psi_1 \# q_0 = q_1$$

**How:** We can regress the vector field with a neural network

$$\mathcal{L}_{FM}(\phi) = \mathbb{E}_{t \sim \mathcal{U}[0,1]} \mathbb{E}_{\mathbf{x} \sim p_t} [\|u_\phi(t, \mathbf{x}) - u(t, \mathbf{x})\|^2]$$

→ We don't know  $u_t(\mathbf{x})$  !

# Conditional Flow Matching

**Key Idea:** Since we have access to samples from  $q_1$ , we can condition the vector field on the target point  $\mathbf{x}_1 \sim q_1$

$$p_t(\mathbf{x}_t) = \int q_1(\mathbf{x}_1) p_{t|1}(\mathbf{x}_t | \mathbf{x}_1) d\mathbf{x}_1 \quad p_{0|1}(\mathbf{x} | \mathbf{x}_1) = \mathcal{N}(0, I)$$
$$\Rightarrow p_{t|1}(\mathbf{x}_t | \mathbf{x}_1) = \mathcal{N}(t\mathbf{x}_1, (1-t)^2 I)$$

The induced  $u_t(\mathbf{x} | \mathbf{x}_1)$  is available in closed form e.g.

$$u_t(\mathbf{x} | \mathbf{x}_1) = \frac{\mathbf{x}_1 - \mathbf{x}}{1 - t}$$

The loss becomes

$$\mathcal{L}_{CFM}(\phi) = \mathbb{E}_{t \sim \mathcal{U}[0,1], \mathbf{x}_1 \sim q_1, \mathbf{x}_t \sim p_t(\mathbf{x} | \mathbf{x}_1)} [\|u_\phi(t, \mathbf{x}) - u_t(\mathbf{x} | \mathbf{x}_1)\|^2]$$

→ Same gradient as  $\mathcal{L}_{FM}$

# FMCPE: Three Main Components

1. **Parameter space flow:** Transport samples from source  $\pi(\theta|y)$  to target  $p(\theta|y)$
2. **Source distribution construction:** Design  $\pi(\theta|y)$  using NPE and learnable kernel

$$\pi(\theta|y) = \int p(\theta|x)q(x|y)dx$$

3. **Joint optimization:** Learn both flows together for improved robustness

# Component 1: Parameter Space Flow

**Goal:** Map samples from source  $\pi(\theta|y)$  to corrected posterior  $p(\theta|y)$

**Flow definition:**

$$\frac{d\theta_t}{dt} = u_\Theta(t, \theta_t, y)$$

**Training objective:** Conditional Flow Matching loss

$$\mathcal{L}_\Theta = \mathbb{E} [\|u_\Theta(t, \theta_t, y) - (\theta_1 - \theta_0)\|^2]$$

where  $\theta_1 \sim p(\theta|y)$  (from calibration) and  $\theta_0 \sim \pi(\theta|y)$  (from source)

## Component 2: Source Distribution

**Key insight:** Use simulation-trained NPE  $p(\theta|x)$  with learned correction

**Source construction:**

$$\pi(\theta|y) = \int p(\theta|x)q(x|y)dx$$

- $p(\theta|x)$ : Pretrained NPE on cheap simulations (frozen)
- $q(x|y)$ : Learnable kernel mapping  $y$  to simulator space

**Why this works:** Bridges the gap between calibration data  $y$  and simulations  $x$

## Component 3: Learning $q(x|y)$

**Approach:** Use Flow Matching to learn  $q(x|y)$

**Flow in simulation space:**

$$\frac{dx_t}{dt} = u_X(t, x_t, y)$$

**Training data generation:** For each calibration pair  $(\theta_i, y_i)$ :

1. Generate  $x_1 = S(\theta_i)$  using the simulator
2. Sample base  $x_0 \sim \mathcal{N}(y, \sigma^2 I)$
3. Train flow to match  $x_0 \rightarrow x_1$

**Result:**  $q(x|y)$  learns to transport from observed space to simulator space

# Joint Training Algorithm

**Key idea:** Train both flows  $u_X$  and  $u_\Theta$  simultaneously

**Training loop:**

**For each minibatch:**

1. Sample  $(\theta_1, y)$  from calibration set  $\mathcal{D}_{\text{cal}}$
2. Generate  $x_1 = S(\theta_1)$  using simulator
3. Sample base points:  $x_0 \sim \mathcal{N}(y, \sigma^2 I)$
4. Map  $x_0$  to simulator space:  $\tilde{x} = T_X(x_0; y)$  using current  $u_X$
5. Sample from source:  $\theta_0 \sim p(\theta|\tilde{x})$  using pretrained NPE
6. Compute flow matching losses for both flows:

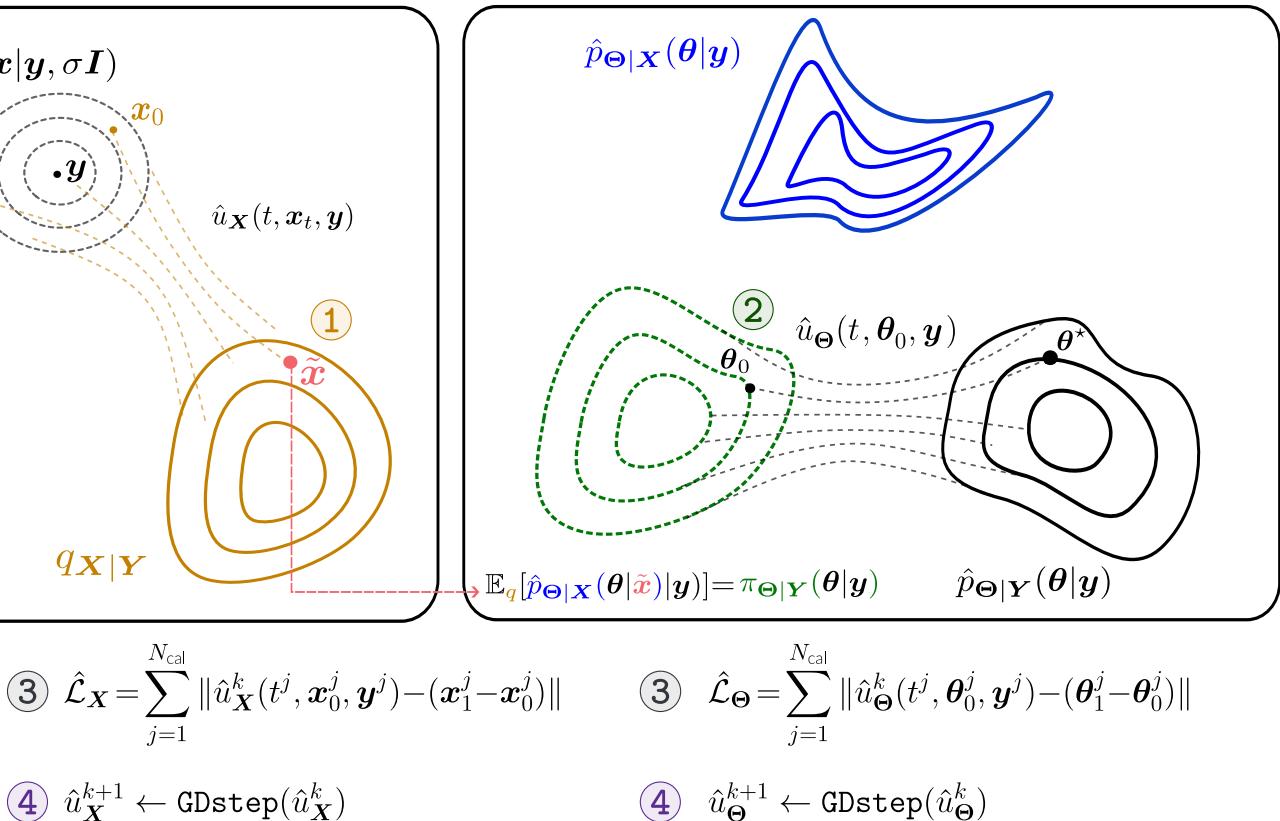
$$\ell_X = \|u_X(t, x_t, y) - (x_1 - x_0)\|^2$$

$$\ell_\Theta = \|u_\Theta(\tau, \theta_\tau, y) - (\theta_1 - \theta_0)\|^2$$

7. Update both  $u_X$  and  $u_\Theta$  using  $\mathcal{L} = \ell_X + \ell_\Theta$

→ Joint training ensures consistency between source and parameter flows

# FMCPE Overview



## Algorithm 1: Iterative update of $\hat{u}_{\Theta}$ and $\hat{u}_X$

**Input:** Calibration set  $\mathcal{D}_{\text{cal}} = \{(\theta^j, \mathbf{y}^j)\}_{1 \leq j \leq N_{\text{cal}}}$ , pre-trained  $p_{\Theta|\mathbf{X}}$   
**for**  $k = 1$  **to**  $K$  **do**

/\*  $\textcircled{1}$  Sample from noise model  
 Sample  $\tilde{\mathbf{x}}$  from  $q_{\mathbf{X}|\mathbf{Y}}(\mathbf{x}|\mathbf{y})$ :

$$\mathbf{x}_0 \sim \mathcal{N}(\mathbf{y}, \sigma^2 \mathbf{I}), \quad \tilde{\mathbf{x}} = \mathbf{x}_0 + \int \underbrace{\hat{u}_X^k(t, \mathbf{x}_0, \mathbf{y})}_{\text{Weights at step } k} dt$$

/\*  $\textcircled{2}$  Sample from proposal  
 Sample  $\theta_0$  from  $\pi_{\Theta|\mathbf{Y}}^k$ :

$$\theta_0 \sim \pi_{\Theta|\mathbf{Y}}^k(\theta|\mathbf{y}) = \mathbb{E}_{q_{\mathbf{X}|\mathbf{Y}}^k} [\hat{p}_{\Theta|\mathbf{X}}(\theta|\tilde{\mathbf{x}})]$$

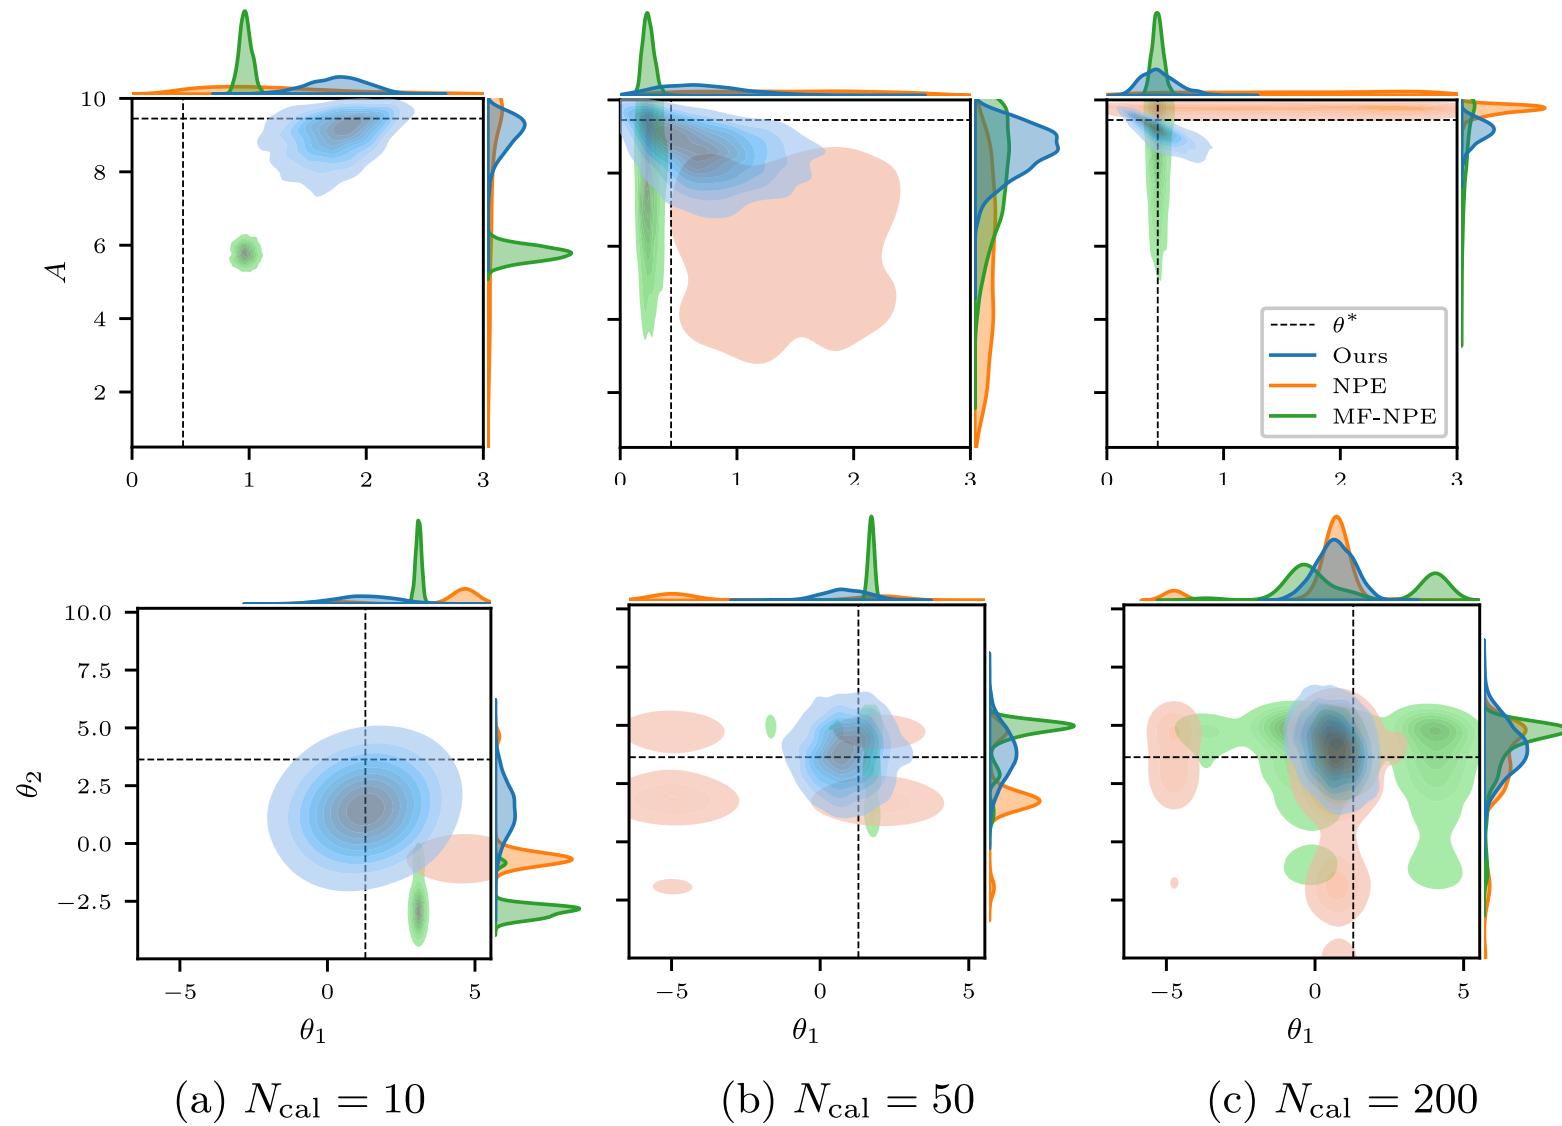
/\*  $\textcircled{3}$  Compute losses with  $\mathbf{x}_1^j = S(\theta^j)$   $\theta_1^j = \theta^j$

$$\hat{\mathcal{L}}_{\Theta} = \sum_{j=1}^{N_{\text{cal}}} \|\hat{u}_{\Theta}^k(t^j, \theta_0^j, \mathbf{y}^j) - (\theta_1^j - \theta_0^j)\| \quad \hat{\mathcal{L}}_X = \sum_{j=1}^{N_{\text{cal}}} \|\hat{u}_X^k(t^j, \mathbf{x}_0^j, \mathbf{y}^j) - (\mathbf{x}_1^j - \mathbf{x}_0^j)\|$$

/\*  $\textcircled{4}$  Gradient update

$$\hat{u}_{\Theta}^{k+1} \leftarrow \text{GDstep}(\hat{u}_{\Theta}^k), \quad \hat{u}_X^{k+1} \leftarrow \text{GDstep}(\hat{u}_X^k)$$

# Results



# SBI Hackathon - Grenoble

**January 21-23, 2026 | Grenoble**

Join us for a hands-on hackathon on the `sbi` package - a simple and powerful tool for simulation-based inference!



<https://github.com/SBI-Grenoble/sbi-hackathon-2026>

**Example: Training an NPE model in just a few lines:**

```
import torch
from sbi.inference import NPE

# define shifted Gaussian simulator
def simulator(theta): return theta + torch.randn_like(theta)

# draw parameters from Gaussian prior
theta = torch.randn(1000, 2)
# simulate data
x = simulator(theta)

# choose sbi method and train
inference = NPE()
inference.append_simulations(theta, x).train()

# do inference given observed data
x_o = torch.ones(2)
posterior = inference.build_posterior()
samples = posterior.sample((1000,), x=x_o)
```

# Thank You!

Questions?

# Appendix

## Common hypothesis

In most settings, the error is assumed to be independent of  $\theta$ :

$$\mathbf{y} \perp \theta \mid \mathbf{x}$$

This is a necessary condition to express  $p(\theta|\mathbf{y})$  using  $p(\theta|\mathbf{x})$  and  $p(\mathbf{x}|\mathbf{y})$ :

$$p(\theta|\mathbf{y}) = \int_{\mathcal{X}} p(\theta|\mathbf{x})p(\mathbf{x}|\mathbf{y})d\mathbf{x}$$

In practice, this is not always true for various physical systems. The low-fidelity simulator might not depend on some of the parameters, using a lower dimensional  $\theta$ .

# NPE Loss Derivation from KL Divergence

We want to minimize the KL divergence for all  $\mathbf{x}$ :

$$\begin{aligned} & \mathbb{E}_{p(\mathbf{x})} [\text{KL}(p(\theta|\mathbf{x}) || q_\phi(\theta|\mathbf{x}))] \\ &= \mathbb{E}_{p(\mathbf{x})} \mathbb{E}_{p(\theta|\mathbf{x})} \left[ \log \frac{p(\theta|\mathbf{x})}{q_\phi(\theta|\mathbf{x})} \right] \end{aligned}$$

Expanding the logarithm:

$$\begin{aligned} &= \mathbb{E}_{p(\mathbf{x})} \mathbb{E}_{p(\theta|\mathbf{x})} [\log p(\theta|\mathbf{x})] \\ &\quad - \mathbb{E}_{p(\mathbf{x})} \mathbb{E}_{p(\theta|\mathbf{x})} [\log q_\phi(\theta|\mathbf{x})] \end{aligned}$$

First term does not depend on  $\phi$ :

$$\min_{\phi} \mathbb{E}_{p(\mathbf{x})} [\text{KL}(p || q_\phi)]$$

Equivalent to maximizing:

$$\max_{\phi} \mathbb{E}_{p(\mathbf{x})} \mathbb{E}_{p(\theta|\mathbf{x})} [\log q_\phi(\theta|\mathbf{x})]$$

Using joint distribution  $p(\theta, \mathbf{x}) = p(\mathbf{x}|\theta)p(\theta)$ :

$$= \max_{\phi} \mathbb{E}_{p(\theta, \mathbf{x})} [\log q_\phi(\theta|\mathbf{x})]$$

Approximate with samples:

$$\mathcal{L}(\phi, \omega) = -\frac{1}{N} \sum_{i=1}^N \log q_\phi(\theta_i | h_\omega(\mathbf{x}_i))$$

where  $(\theta_i, \mathbf{x}_i) \sim p(\theta)p(\mathbf{x}|\theta)$